

Multilayer perceptrons

Recall from last two lectures

- Least squares classifier:

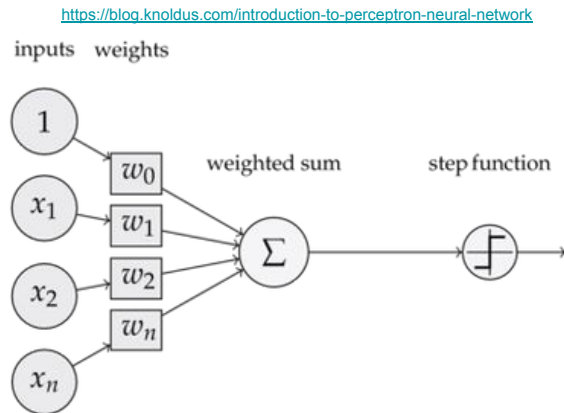
$$\hat{\mathbf{y}} = \mathbf{W}\mathbf{x} \quad L(\mathbf{y}, \hat{\mathbf{y}}) = \frac{1}{2} \sum_{i=1}^m (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2$$

- Perceptron:

$$\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x}) \quad L(y, \hat{y}) = - \sum_{i=1}^m \mathbf{w}^\top \mathbf{x}_i y_i$$

- Logistic regression:

$$\hat{y} = \sigma(\mathbf{w}^\top \mathbf{x}) \quad L(y, \hat{y}) = - \sum_{i=1}^m y_i \log \sigma(\mathbf{w}^\top \mathbf{x}) + (1 - y_i) \log(1 - \sigma(\mathbf{w}^\top \mathbf{x}))$$



Fully connected NNs

- In DL, you will usually see this form:


$$\mathbf{h} = g(\mathbf{W}\mathbf{x} + \mathbf{b})$$

W: Matrix of weights, one vector per neuron

x: One input example (vector)

b: Vector of biases, one scalar per neuron

h: Hidden layer response




Hyperparameters

- Number of layers
- Propagation types: fully connected, convolutional (later)
- Activation function(s)
- Loss function(s) and parameters (also later)
- Training iterations and batch sizes (some of it today)


Input layer

- A vectorized version of the input data
- Sometimes preprocessed
- Weights connect to the hidden layer




Hidden layer(s)

- Number of hidden layers
- Number of neurons
- Topology - expanding or bottleneck
- Strongly application-dependent
- Currently no optimal way to design in advance



Output layer

- For regression:
 - Linear outputs with MSE
- For classification
 - Softmax units (logistic for two classes)
- - and many, many more...



Activation functions


$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$

$$\tanh'(x) = 1 - \tanh(x)^2$$

$$\text{relu}'(x) = \text{step}(x)$$


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$



ReLU

$$\max(0, x)$$




The multilayer perceptron

- A one-way computational chain
- First, we take the input \mathbf{x} and process it:
$$\mathbf{h}_1 = g_1(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1)$$
- This gives a first *hidden* representation. Next layer:
$$\mathbf{h}_2 = g_2(\mathbf{W}_2\mathbf{h}_1 + \mathbf{b}_2)$$
- - and so on

Universal approximation

- Think of MLPs as a big function block with many free parameters
- Even with a single hidden layer, *any function* can be represented
 - - as long as we use a non-linearity
- In practice, single-layer nets may not be trained well to a task
- Instead, we *go deep* and reduce the number of neurons per layer



Recall maximum likelihood



- We setup a (neural) model p , which tries to predict an output (label) from an input (e.g. image)
- The ML estimate of its parameters is done using a *training set*:
$$W_{\text{ML}} = \arg \max E_{\mathbf{x} \sim \hat{p}_{\text{data}}} \log p(y|\mathbf{x})$$
- where the expectation is simply the mean over the m training examples

Generalization error of $W_{\text{ML}} = \arg \max_{\mathbf{x} \sim \hat{p}_{\text{data}}} E \log p(y|\mathbf{x})$

- The problem
 - We do not have access to the full population
 - Instead, we get a limited sample, the *empirical distribution* \hat{p}_{data}
- Now we want our model p to follow this empirical distribution - sort of
 - We actually want our model to predict future *test* cases
 - In the end, what we *really* want is high test classification accuracy, but we'll have to do with something we can optimize upon, cross-entropy on the training set




Over- and underfitting

- Take an example
- Best fit line
- Best value of w with associated losses



Over- and underfitting

- Another example (training points are sampled on a 2nd degree polynomial)
- Linear regression “fails”
- One more parameter is appropriate
- **This also happens for more realistic training data!**




Train-val-test split

- Usually we get a training set for learning
- The test set is not to be used during learning
 - Sometimes hidden behind an evaluation server
- We can hold out a *validation set* for tuning
 - Also sometimes given by the benchmark
- This allows for design of e.g. network topology



Network capacity

- The common method to increase capacity of an underfitting model is more weights (more neurons)
- Now we can evaluate this hyperparameter:
 - Full training and evaluation on test set
 - Use the validation set



Exercise

- Introduction to PyTorch

<https://drive.google.com/file/d/1IMC9OJjr-MsgiLTCpAsBNY1JWmXf-cVn/view?usp=sharing>

- Multi Layer Perceptron

<https://drive.google.com/file/d/1Go41K2FcH5Ng3sLZKyNm-q55rXCqXJyd/view?usp=sharing>

Introduction to PyTorch

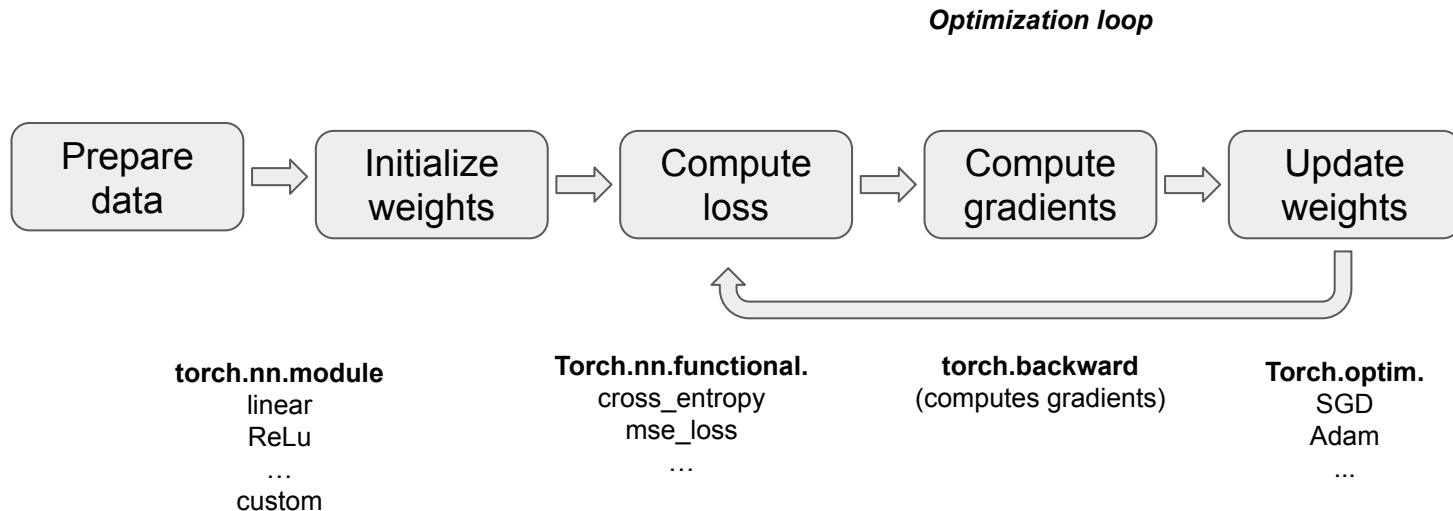
- Introduced by FAIR in 2017
- Built with Python as a primary programming language
- Modular and flexible
- Short learning curve & easy to debug

NumPy v/s PyTorch

- Both API's are very similar
- NumPy arrays -> PyTorch tensors
 - GPU support
 - Automatic differentiation
- Some other differences:
 - NumPy "axis" -> PyTorch "dim"
 - NumPy "reshape" -> PyTorch "view"

Last week - Linear classifiers

- $y = ax + b$
- Using PyTorch



Challenge

- Build a Multi Layer Perceptron (MLP) with one ReLU-activated hidden layer
- Use CIFAR10 dataset
- Use the validation set to determine number of hidden neurons
- Use the test set to get an unbiased estimate of your model's performance on the real data distribution