

Regularization

What is regularization?

- One definition in the book
 - Any modification to the learning that aims to reduce generalization error, but not training error
- The problem
 - With high-capacity models (many layers, many parameters), there is a high risk of overfitting
 - If we just naively reduce the model capacity, we risk underfitting
 - Hitting the “sweet spot” in the middle is difficult
- The solutions
 - Regularization allows you to use high-capacity models
 - Some constraint or clever strategy *regularizes* this model to behave more smoothly
 - We are thus able to trade-off between capacity and robustness

Examples of regularization

- Parameter norm penalties
 - Classical problem in optimization: unstable parameter solutions
- Data augmentation
 - Artificially expand the training dataset
 - Allows the model to see a larger variation of training examples
- Noise injection
 - Add noise to parts of the model
 - Input noise: a bit like dataset augmentation
 - Noise on the units: dropout

Recall maximum likelihood - again

- We set up a model to predict y given \mathbf{x} :
 $p(y|\mathbf{x})$
- Now, we change notation a bit, realizing that y is really also a function of our parameters, which we now call θ :
 $p(y|\mathbf{x}, \theta)$
- In ML estimation, what we really had was a likelihood function of the whole training data sample (X, Y) :
$$\theta_{\text{ML}} = \arg \max p(Y|X, \theta) = \arg \max \prod_i p(y_i|\mathbf{x}_i, \theta)$$

$$\theta_{\text{ML}} = \arg \max p(Y|X, \theta) = \arg \max \prod_i p(y_i | \mathbf{x}_i, \theta)$$

MAP inference

- Now let's think of the problem another way: $p(\theta|X, Y)$
- Whenever there's a conditional, we may think of good ol' Bayes
- If we use Bayes' theorem, this part is equal to (just pretend that A is θ and that B is X, Y):

$$p(\theta|X, Y) = \frac{p(X, Y|\theta)p(\theta)}{p(X, Y)}$$

- Since we are trying to maximize, we can neglect the denominator:

$$p(\theta|X, Y) \propto p(X, Y|\theta)p(\theta)$$

- What about $p(\theta)$?

$$p(\theta|X, Y) \propto p(X, Y|\theta)p(\theta)$$

MAP inference

- If we let $p(\theta) = \text{const.}$, we are back to something familiar:

$$p(X, Y|\theta) = \frac{p(X, Y, \theta)}{p(\theta)} = \frac{p(Y|X, \theta)p(X|\theta)p(\theta)}{p(\theta)} = p(Y|X, \theta)p(X|\theta)$$


- If not, we have what is called a *prior* on our parameters:

$$p(\theta|X, Y) \propto p(X, Y|\theta)p(\theta)$$

- The left part that models θ is the *posterior* distribution

- The solution to this is called a *MAP estimator*:

$$\theta_{\text{MAP}} = \arg \max p(Y|X, \theta)p(\theta)$$


$$\theta_{\text{ML}} = \arg \max p(Y|X, \theta)$$

MAP inference

- Let's look at a known case
 - Gaussian output distribution with a linear regression function: $y \sim \mathcal{N}(f(\mathbf{x}), \sigma^2)$
 - But now, let's also assume a standard Gaussian prior on θ : $\theta \sim \mathcal{N}(0, \mathbf{I})$
- We then get the following:

$$\begin{aligned}\theta_{\text{MAP}} &= \arg \max_{\theta} p(Y|X, \theta)p(\theta) \\ &= \arg \max_{\theta} \left(\prod_{i=1}^m \mathcal{N}(y_i; f(\mathbf{x}_i), \sigma^2) \right) \mathcal{N}(\theta; 0, \mathbf{I}) \\ &= \arg \max_{\theta} \mathcal{N}(\theta; 0, \mathbf{I}) \prod_{i=1}^m \mathcal{N}(y_i; f(\mathbf{x}_i), \sigma^2)\end{aligned}$$

$$\theta_{\text{MAP}} = \arg \max \mathcal{N}(\theta; 0, \mathbf{I}) \prod_{i=1}^m \mathcal{N}(y_i; f(\mathbf{x}_i), \sigma^2)$$

MAP inference

- If we plug in the Gaussian pdf and remove constants, we get:

$$\theta_{\text{MAP}} = \arg \max e^{-\frac{1}{2}(\theta-0)^2} \prod_i e^{-\frac{1}{2}(y_i - \hat{y}_i)^2} \quad \hat{y} = f(\mathbf{x})$$

- As usual, we take the negative logarithm to arrive at a loss to minimize:

$$\theta_{\text{MAP}} = \arg \min \left(\frac{1}{2} \|\theta\|^2 + \frac{1}{2} \sum_i (y_i - \hat{y}_i)^2 \right)$$

- What if the “likelihood part” (rightmost) is multinoulli?

- No problem - the likelihood and the prior are separated in log-space:

$$\theta_{\text{MAP}} = \arg \min \left(\frac{1}{2} \|\theta\|^2 - \sum_i y_i \log \hat{y}_i \right)$$

Weight decay

- All this leads to a modified loss function that takes into account the prior:

$$L(y, \hat{y}; \theta) = L(y, \hat{y}) + \alpha \Omega(\theta)$$

- When using a Gaussian prior:

$$\Omega(\theta) = \frac{1}{2} \|\theta\|^2$$

New hyperparameter



- we also call it
 - L_2 regularization/penalty
 - Tikhonov regularization
 - Ridge regression (mostly reserved for regression models)
- In deep learning: *weight decay*

$$L(y, \hat{y}; \theta) = L(y, \hat{y}) + \alpha \Omega(\theta)$$

Norm penalties

- Another common prior is the *Laplacian*:

$$\theta \sim \text{Laplace}(0, 1) = \frac{1}{2} e^{-|\theta|}$$

- which leads to the L_1 penalty:

$$\Omega(\theta) = \|\theta\|_1$$

$$L(y, \hat{y}; \theta) = L(y, \hat{y}) + \alpha \Omega(\theta)$$

Norm penalties

- Computing gradients is often very easy
- The first loss term we already solved with backprop
- The regularization term has a simple solution for L_1 and L_2
 - L_1 : $\nabla_{\theta} \Omega = \frac{\theta}{|\theta|} = \text{sign}(\theta)$
 - L_2 : $\nabla_{\theta} \Omega = \theta$

Weight decay

- Let's look at the L_2 case

$$L(y, \hat{y}; \theta) = L(y, \hat{y}) + \frac{1}{2}\alpha\|\theta\|^2$$

- with the gradient

$$\nabla_{\theta} L(y, \hat{y}) + \alpha\theta$$



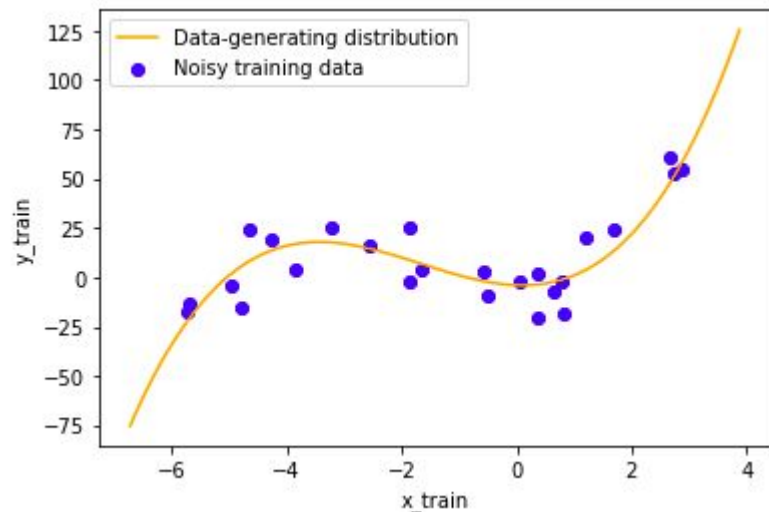
- When we take small steps along the negative:

$$\theta \leftarrow \theta - \epsilon (\nabla_{\theta} L(y, \hat{y}) + \alpha\theta)$$

- you should be able to see how this “shrinks” or “decays” the weights


Weight decay

- Simple example




Augmentation

- When training data is scarce, we can artificially expand the training set
- Typical strategies
 - Noise
 - Additive, e.g. Gaussian
 - Multiplicative, e.g. Bernoulli (Dropout)
 - Affine transformation
 - Translations
 - Rotations
 - Scaling
 - Shearing
 - Truncation
 - Cropping
 - Non-linear operations
 - Brightness



<https://hazyresearch.github.io/snorkel/blog/tanda.html>

Early stopping



Early stopping

Number of updates per epoch

Best validation performance so far

New best performance achieved

Reset patience

Update parameters and performance

Algorithm 7.1 The early stopping meta-algorithm for determining the best amount of time to train. This meta-algorithm is a general strategy that works well with a variety of training algorithms and ways of quantifying error on the validation set.

Let n be the number of steps between evaluations.

Let p be the “patience,” the number of times to observe worsening validation set error before giving up.

Let θ_o be the initial parameters.

$\theta \leftarrow \theta_o$

$i \leftarrow 0$

$j \leftarrow 0$

$v \leftarrow \infty$

$\theta^* \leftarrow \theta$

$i^* \leftarrow i$

while $j < p$ **do**

 Update θ by running the training algorithm for n steps.

$i \leftarrow i + n$

$v' \leftarrow \text{ValidationSetError}(\theta)$

if $v' < v$ **then**

$j \leftarrow 0$

$\theta^* \leftarrow \theta$

$i^* \leftarrow i$

$v \leftarrow v'$

else

$j \leftarrow j + 1$

end if

end while

Best parameters are θ^* , best number of training steps is i^* .

Dropout



- A method for regularizing the network during training
- First the method:
 - During forward pass, randomly set neurons to zero with a probability $p < 1$
 - This “weakens” the network so that it only has a “power” of $1 - p$
 - To compensate, multiply all other neurons by $1 / (1 - p)$
 - The backward pass is trivial: some responses are just zero and provide no gradient
- Common values for p
 - Input units: between 0 and 0.2
 - Hidden units: 0.5
 - Output units (e.g. linear, softmax): **always 0!**

Dropout

- Slightly more formal:
 - Multiply each neuron by a Bernoulli random variable $d \sim \text{Bernoulli}(p)$
 - The rescaling of $1 / (1 - p)$ ensures that the net input to any neuron stays the same
 - Each pass through a layer looks like this (no rescaling):
$$\mathbf{h}_i = \phi(\mathbf{W}_i \cdot (\mathbf{h}_{i-1} \odot \mathbf{d}_{i-1}) + \mathbf{b}_i)$$
 - And with rescaling:
$$\mathbf{h}_i = \phi\left(\mathbf{W}_i \cdot \left(\mathbf{h}_{i-1} \odot \frac{\mathbf{d}_{i-1}}{1-p}\right) + \mathbf{b}_i\right)$$


Dropout performance

- Paper: Srivastava, Nitish, et al. "Dropout: a simple way to prevent neural networks from overfitting." *JMLR*'14.



Dropout subnet interpretation

- Each time you drop a set of units, you effectively create a new “subnetwork”
- These subnetworks obviously share the same weights
- As such, repeated training like this, effectively corresponds to a variant of *bagging*
 - At each minibatch, dropout samples a new model
 - This model sees a small, random subset of the data
 - In the end, this “ensemble” of models is combined



Exercise

- You are given only 1000 MNIST training examples
- The baseline model overfits (100% train acc, ~89% valid acc)
- Can you improve this with regularization?

Bonus:

- What performance can you achieve using the whole dataset?