


# CNNs


# Definition of CNNs

- One or more layers in the net use convolutions




# Convolutions

- Confusing terminology
  - Convnets actually perform *cross-correlation*
- Computing a single output pixel:  
$$S(i, j) = \sum_m \sum_n I(i + m, j + n)K(m, n)$$




# Convolutions


$$S(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n)$$




## Full connectivity



## Sparse connectivity




# Receptive field sizes





# Receptive field sizes

- As we go through higher layers, individual neurons “see” more of the input
- Add to this the increasing number of non-linearities
- This way, higher layers learn more abstract concepts



# Weight matrices as filters

- Early layers in a CNN learn to extract primitive features
- 96 of the learned AlexNet filters (11x11x3) are shown here




Krizhevsky et al., 2012




# Higher-level abstractions

- Girshick et al., 2014




# Padding

- **Input**
- **Filter**
- **Output feature map**
- Helps maintaining dimensions in the output feature map
- Kernel size:
  - 3
- Padding:
  - 1
- Stride:
  - 1



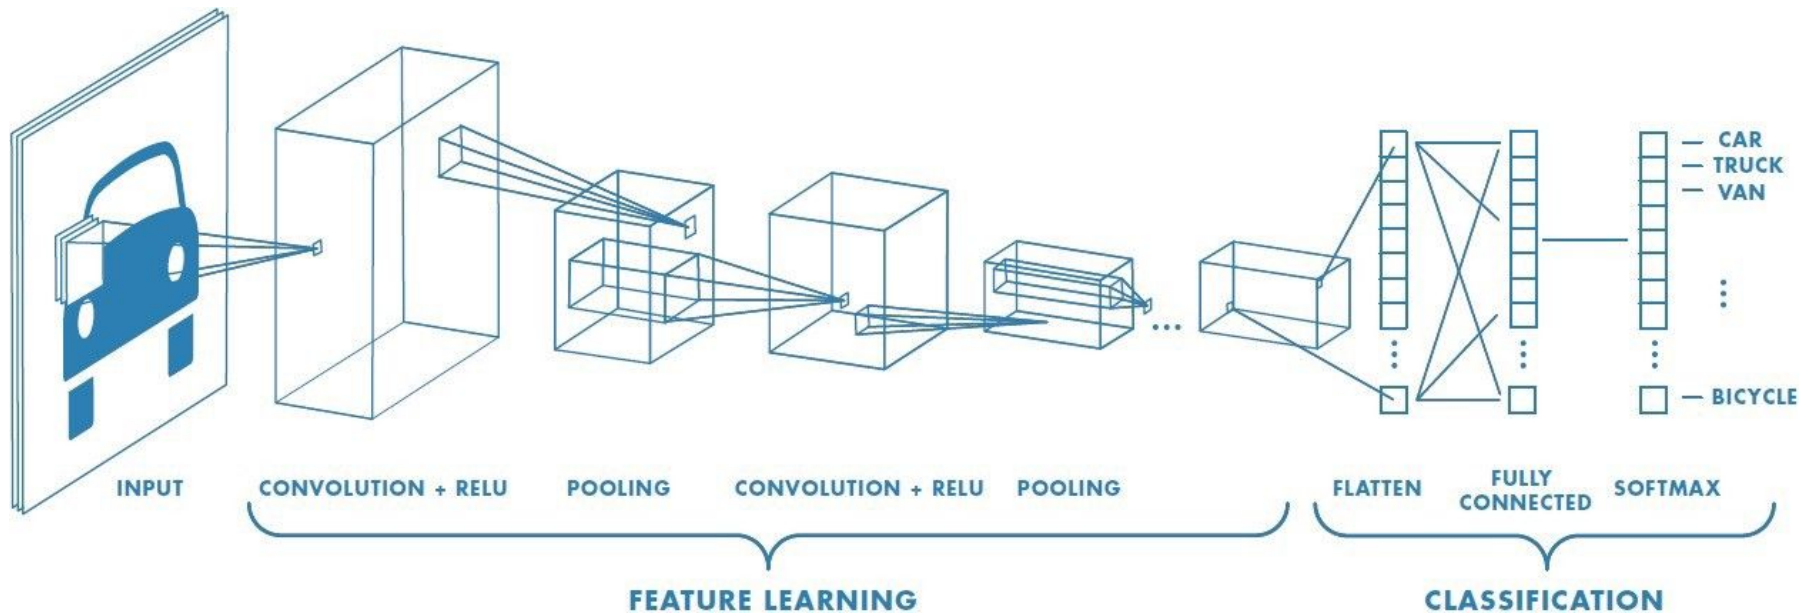
# Strided convolutions

- **Input**
- **Filter**
- **Output feature map**
- Reduces data
- Maintains local invariance
- Kernel size:
  - 3
- Padding:
  - 1
- Stride:
  - 2





# Pooling

- Data reduction
- Local invariance




# Pooling (1D)



# Strided pooling



- Similar to strided convolutions, but no parameters
- The stride (here 2) controls how much to downsample
- The pool size controls the “locality”





# Pooling in 2D


- Here we use a stride of 2 and a pool size of 2x2




Some famous CNNs



# LeNet




# AlexNet



# VGGNet

- Simonyan & Zisserman, 2015

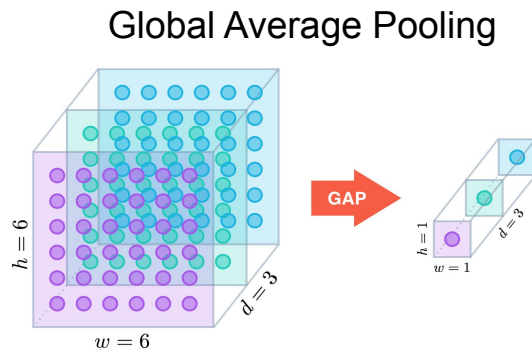


ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 <b>LRN</b>	conv3-64 <b>conv3-64</b>	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 <b>conv3-128</b>	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 <b>conv1-256</b>	conv3-256 conv3-256 <b>conv3-256</b>	conv3-256 conv3-256 conv3-256 <b>conv3-256</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

# The all convolutional net

- Springenberg et al., 2015
- Replace all pooling operations with stride-2 convolutions
- Use *global average pooling* in the output layer

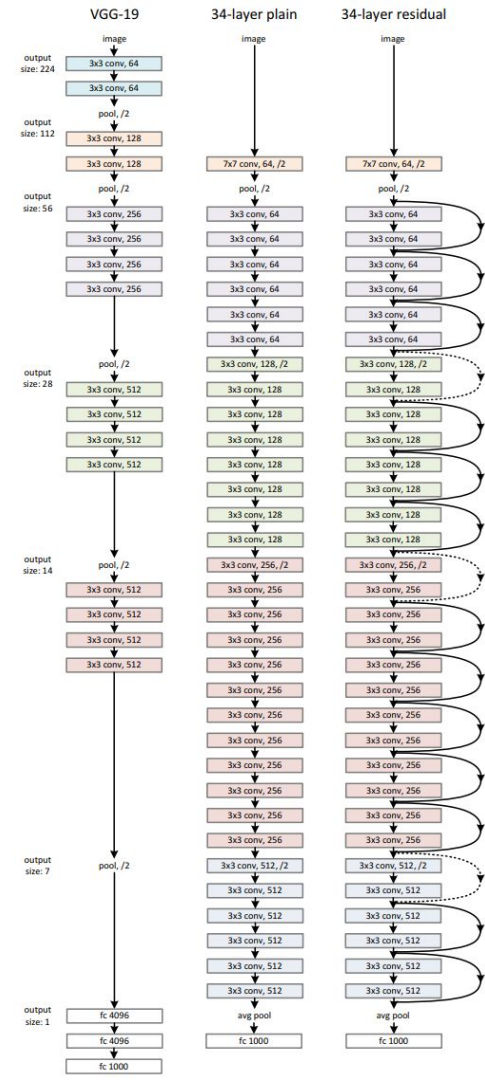



ImageNet model	
Layer name	Layer description
input	Input $224 \times 224$ RGB image
conv1	$11 \times 11$ conv. 96 ReLU units, stride 4
conv2	$1 \times 1$ conv. 96 ReLU, stride 1
conv3	$3 \times 3$ conv. 96 ReLU, stride 2
conv4	$5 \times 5$ conv. 256 ReLU, stride 1
conv5	$1 \times 1$ conv. 256 ReLU, stride 1
conv6	$3 \times 3$ conv. 256 ReLU, stride 2
conv7	$3 \times 3$ conv. 384 ReLU, stride 1
conv8	$1 \times 1$ conv. 384 ReLU, stride 1
conv9	$3 \times 3$ conv. 384 ReLU, stride 2, dropout 50 %
conv10	$3 \times 3$ conv. 1024 ReLU, stride 1
conv11	$1 \times 1$ conv. 1024 ReLU, stride 1
conv12	$1 \times 1$ conv. 1000 ReLU, stride 1
global_pool	global average pooling ( $6 \times 6$ )
softmax	1000-way softmax



<https://towardsdatascience.com/iclr-2015-striving-for-simplicity-the-all-convolutional-net-with-interactive-code-manual-b4976e206760>



# ResNet

- He. et al., 2015
- Residual “shortcuts” allow for gradient flow
- 100-layer barrier overcome for the first time




# Densely connected CNNs

- Huang et al., 2016
- In dense blocks, all feature maps are concatenated
- Downsampling only occurs between blocks



**Figure 3:** Comparison of the DenseNets and ResNets top-1 error rates (single-crop testing) on the ImageNet validation dataset as a function of learned parameters (*left*) and FLOPs during test-time (*right*).



# Challenge - CIFAR10

Previously: KNN (30-40 %), MLP (50-60 %)

- Design and train a CNN on CIFAR10 (> 80 %)
  - Increase the channels while decreasing the resolution through pooling or strided convolutions
  - Potential help: ~ 4 conv layers, use Adam, batchnorm, data augmentation, batchsize: 128, num\_workers: 4
- Bonus: Transfer Learning (> 90 %)
  - Load a [pretrained resnet 18](#) model and replace the head with a new head to classify CIFAR10 images
  - Potential help:
    - resnet18.fc = [your new head]
    - [Upscale](#) the CIFAR10 images to the image size it was pretrained on (224 x 224)
    - Lower the learning rate to avoid destroying the pretrained weights

airplane



automobile



bird



cat



deer



dog




frog



horse



ship



truck





# CNN topologies (classification)

- Notice how many CNNs follow a similar pattern

