


# Autoencoders

# What is an autoencoder?

- Here's one view:
  - In comes an input example  $\mathbf{x}$
  - Then the *encoder*  $f$  is applied
  - The output  $\mathbf{h}$  of the encoder is called a *code* or a *latent representation/code*
  - The *decoder*  $g$  tries to reproduce  $\mathbf{x}$  using  $\mathbf{h}$



# What is an autoencoder?

- More formally:  
 $\mathbf{r} = g(\mathbf{h}) = g(f(\mathbf{x}))$
- If we're doing a good job, hopefully we'll have:  
 $\mathbf{r} \approx \mathbf{x}$

- In general, we get the loss for AEs:


$$L(\mathbf{x}, g(f(\mathbf{x}))) = L(\mathbf{x}, \mathbf{r})$$

- Predicting black-white pixels (MNIST):

$$L(x, r) = -\frac{1}{m} \sum_{i=1}^m x \log r \quad \text{Sigmoidal outputs}$$


- Predicting real-valued pixels (general images):

$$L(x, r) = \frac{1}{2m} \sum_{i=1}^m (x - r)^2 \quad \text{Linear outputs}$$




# What is an autoencoder?

- *Overcomplete* version
- More components in the latent code than in the input code
- Careful with this one - potentially difficult to train
  - Identity mapping:  $g(f(\mathbf{x}))$  just becomes the identity function
- Solution:
  - Regularization
  - Weird tricks (later)




(b) Shallow overcomplete




(d) Deep overcomplete

# What is an autoencoder?

- *Undercomplete* version
- In general much more interesting:
  - A compressor of data
  - A non-linear PCA machine
  - An algorithm *forced* to learn what's “interesting” in our data







(a) Shallow undercomplete



(c) Deep undercomplete

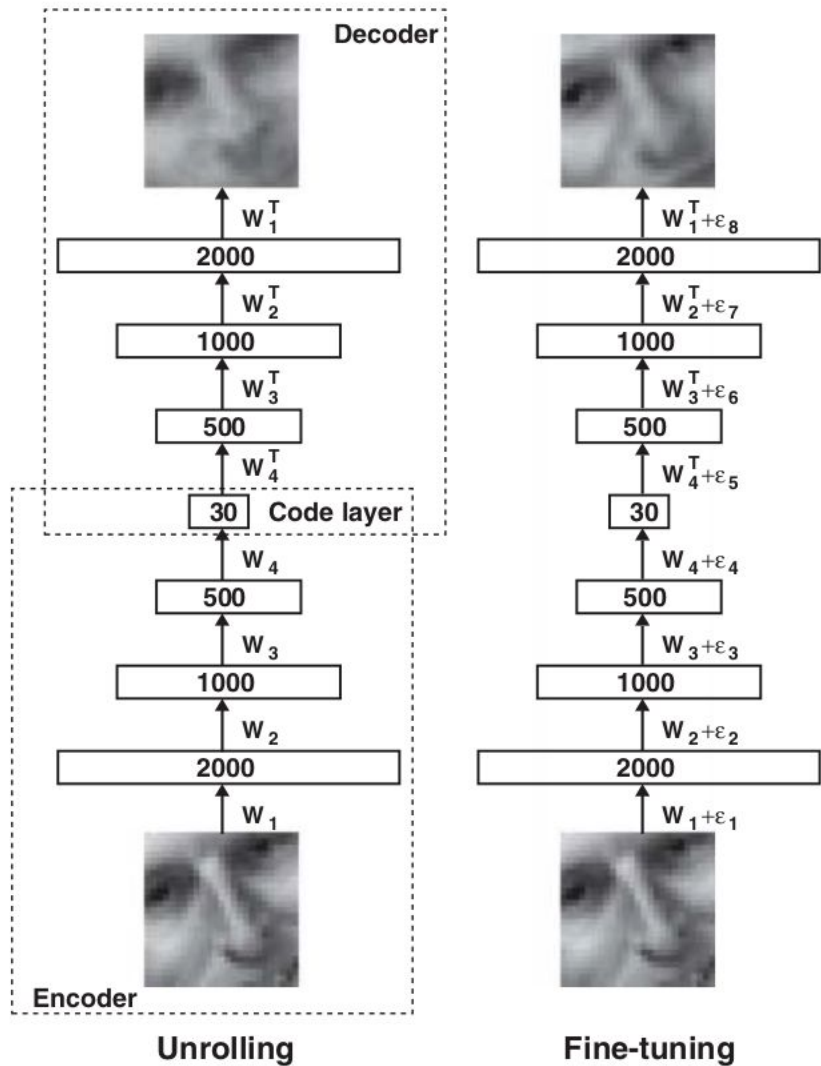
# The rise of AEs

- Before ReLU, deep nets were seen as untrainable
- CNNs were working with some depth, but also quite few feature maps
- Hinton & Salakhutdinov came up with a solution in 2006
- First step:
  - Greedy layer-wise pretraining
  - Each double-arrow means  $f$  going upwards and  $g$  going downwards
  - RBMs are just a special type of autoencoders




# The rise of AEs

- Upon completion of pretraining, the weights are “unrolled” and the whole net is jointly trained further (fine-tuning)
  - This made deep nets (here 8 layers) actually *work* with sigmoids
- Later, Hinton himself used ReLUs, requiring no pretraining
- The method of pretrained nets remained popular in another form: CNN backbones



# Codes learned by AEs

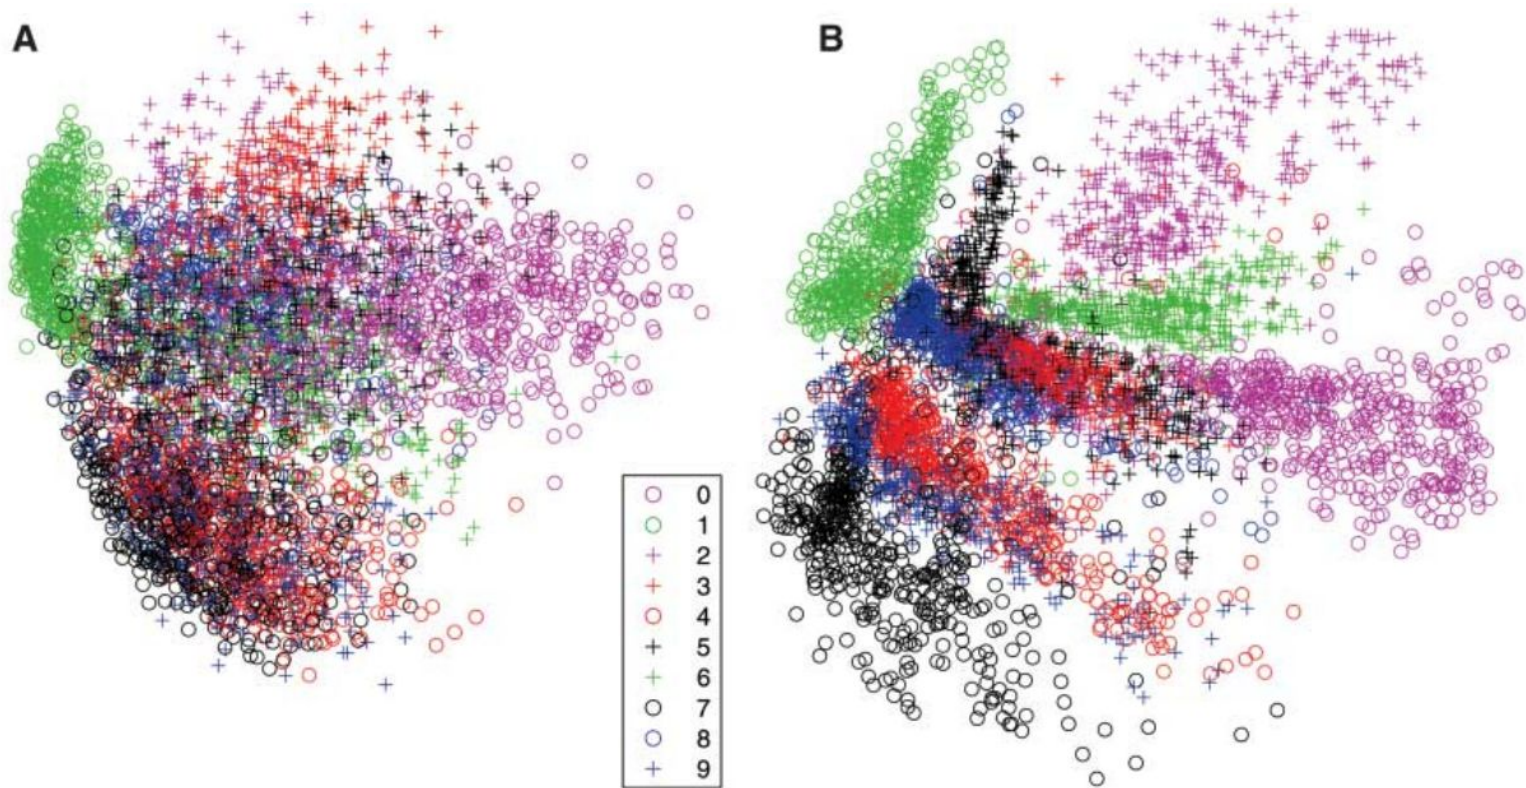
**Fig. 2.** (A) Top to bottom: Random samples of curves from the test data set; reconstructions produced by the six-dimensional deep autoencoder; reconstructions by “logistic PCA” (8) using six components; reconstructions by logistic PCA and standard PCA using 18 components. The average squared error per image for the last four rows is 1.44, 7.64, 2.45, 5.90. (B) Top to bottom: A random test image from each class; reconstructions by the 30-dimensional autoencoder; reconstructions by 30-dimensional logistic PCA and standard PCA. The average squared errors for the last three rows are 3.00, 8.01, and 13.87. (C) Top to bottom: Random samples from the test data set; reconstructions by the 30-dimensional autoencoder; reconstructions by 30-dimensional PCA. The average squared errors are 126 and 135.





# 2D codes learned by an AE on MNIST

**Fig. 3.** (A) The two-dimensional codes for 500 digits of each class produced by taking the first two principal components of all 60,000 training images. (B) The two-dimensional codes found by a 784-1000-500-250-2 autoencoder. For an alternative visualization, see (8).



# Another use of AEs

- Doumanoglou et al., CVPR'16

([https://www.cv-foundation.org/openaccess/content\\_cvpr\\_2016/papers/Doumanoglou\\_Recovering\\_6D\\_Object\\_CVPR\\_2016\\_paper.pdf](https://www.cv-foundation.org/openaccess/content_cvpr_2016/papers/Doumanoglou_Recovering_6D_Object_CVPR_2016_paper.pdf))




Figure 2: Framework Overview. After patch extraction, RGBD channels are given as input to the Sparse Autoencoder. The annotation along with the produced features of the middle layer are given to a Hough Forest, and the final hypotheses are generated as the modes of the Hough voting space. After refining the hypotheses using joint registration, we estimate the next-best-view using a pose-to-lead mapping learnt from the trained Hough Forest.




# SAEs

- The section on sparse AEs in the book presents an alternative view on MAP inference, when used in this context
- What it boils down to is a prior on the hidden  $\mathbf{h}$  instead of the normal prior on our weights  $\mathbf{W}$


# SAEs

- Sparse AEs introduce a prior on the hiddens, formalized as:  
 $L(x, r) + \Omega(\mathbf{h})$
- The most common use is  $L_1$ :  
 $L(x, r) + \lambda \sum_i |h_i|$



# DAEs

- Denoising AEs are trained by corrupting the input
  - Denote the corrupted input  $\tilde{\mathbf{x}}$
- The output is kept “clean”
- The loss looks like this:  
 $L(\mathbf{x}, g(f(\tilde{\mathbf{x}})))$



# DAEs

---

## Extracting and Composing Robust Features with Denoising Autoencoders

---

Pascal Vincent  
Hugo Larochelle

Yoshua Bengio

Pierre-Antoine Manzagol

Université de Montréal, Dept. IRO, CP 6128, Succ. Centre-Ville, Montréal, Québec, H3C 3J7, Canada

VINCENTP@IRO.UMONTREAL.CA

LAROCHEH@IRO.UMONTREAL.CA

BENGIOY@IRO.UMONTREAL.CA

MANZAGOP@IRO.UMONTREAL.CA


### Abstract

Previous work has shown that the difficulties in learning deep generative or discriminative models can be overcome by an ini-

to ponder the difficult problem of inference in deep directed graphical models, due to “explaining away”. Also looking back at the history of multi-layer neural networks, their difficult optimization (Bengio et al., 2007; Bengio, 2007) has long prevented reaping the ex-

# DAEs


- Noise regularizes the net and produces more smooth filters




(a) No destroyed inputs

(b) 25% destruction

(c) 50% destruction



(d) Neuron A (0%, 10%, 20%, 50% destruction)




(e) Neuron B (0%, 10%, 20%, 50% destruction)



# DAEs

- DAEs are better at learning the structure or the *manifold* of  $p_{\text{data}}$



**Figure 2. Manifold learning perspective.** Suppose training data ( $\times$ ) concentrate near a low-dimensional manifold. Corrupted examples ( $\bullet$ ) obtained by applying corruption process  $q_{\mathcal{D}}(\tilde{X}|X)$  will lie farther from the manifold. The model learns with  $p(X|\tilde{X})$  to “project them back” onto the manifold. Intermediate representation  $Y$  can be interpreted as a coordinate system for points on the manifold.



# DAEs

- Let's look at an example:

[https://colab.research.google.com/drive/1jXQaleSqksOrWhDPtze6YnieuuAc\\_mAe?usp=sharing](https://colab.research.google.com/drive/1jXQaleSqksOrWhDPtze6YnieuuAc_mAe?usp=sharing)

# CAEs

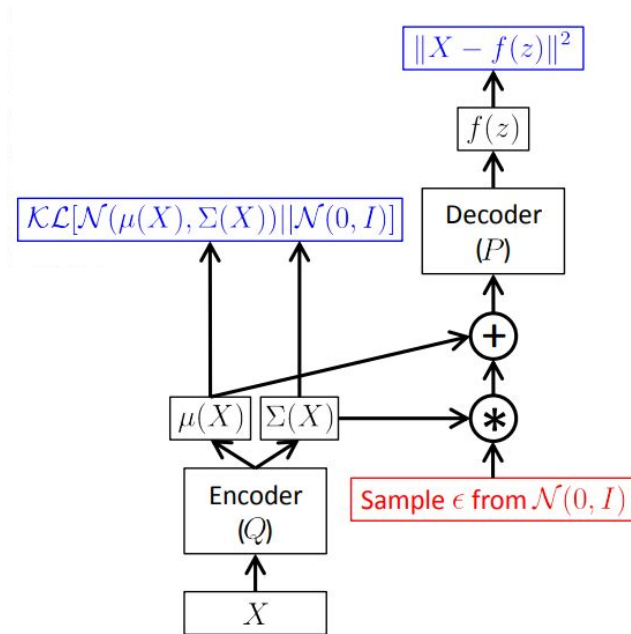
- Contractive AEs try to achieve insensitivity to input perturbations directly on the hidden code
  - Remember that DAEs tried to achieve this on the reconstruction directly
- The term is now defined as a penalty on the “speed” of change of  $\mathbf{h}$ :

$$\Omega(\mathbf{h}) = \lambda \left\| \frac{\partial \mathbf{h}}{\partial \mathbf{x}} \right\|_F^2 = \lambda \left\| \frac{\partial f(\mathbf{x})}{\partial \mathbf{x}} \right\|_F^2$$

- This has the effect of “contracting” nearby input examples to nearby hidden codes

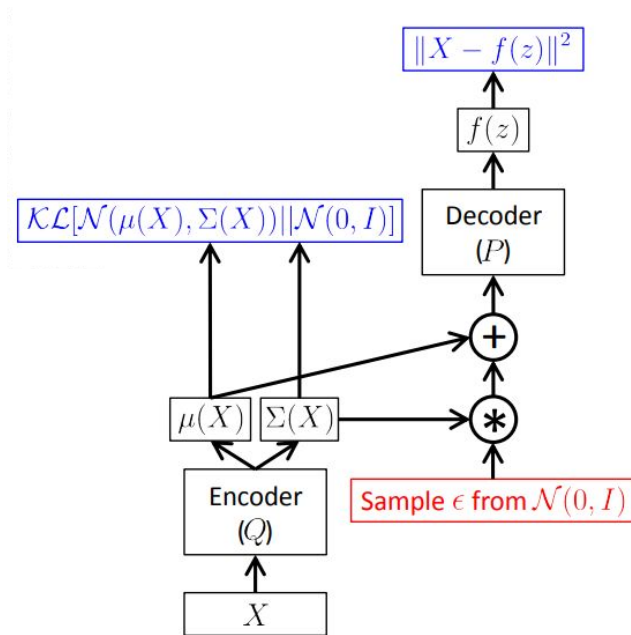
# VAEs

- Variational AEs are trained to produce a latent mean and variance



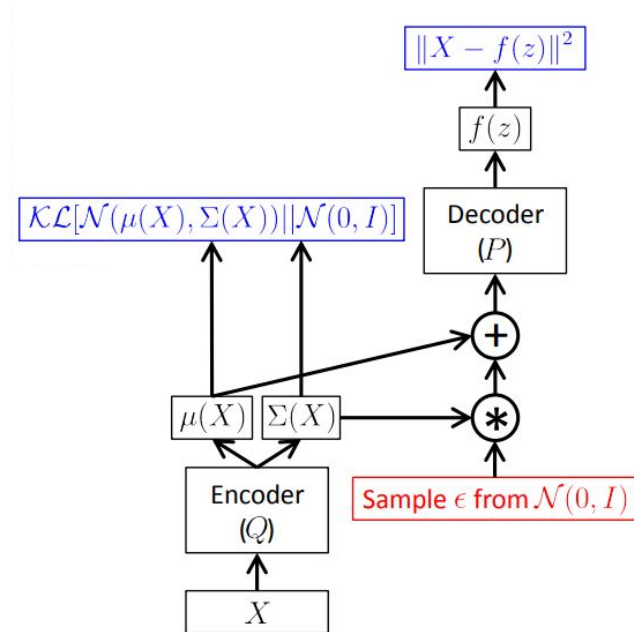
# VAEs

- VAEs then use a statistical *divergence* to make the latent code Gaussian



# VAEs

- When properly trained:
  - The reconstruction is good at reproducing the input
  - The latent mean/variance are Gaussian
- **This makes sampling easy!**



# VAEs


- Here's the principle of the added KL-based regularization:



<https://towardsdatascience.com/understanding-variational-autoencoders-vaes-f70510919f73>

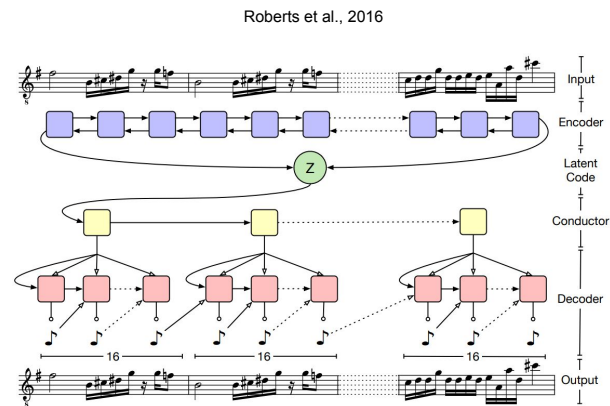

# VAEs

- This allows for easy sampling and interpolation:



<https://towardsdatascience.com/understanding-variational-autoencoders-vaes-f70510919f73>

# VAEs




<https://towardsdatascience.com/intuitively-understanding-variational-autoencoders-1bfe67eb5daf>



# Convolutional AEs


- Finally, you can of course use CNNs for doing better autoencoding of images
- The key ingredient is *transposed convolution*
  - Early literature misused the term *deconvolution* for the same principle



[https://www.researchgate.net/figure/Illustration-of-convolutional-autoencoder-In-this-example-the-CAE-comprises-two\\_fig3\\_306081538](https://www.researchgate.net/figure/Illustration-of-convolutional-autoencoder-In-this-example-the-CAE-comprises-two_fig3_306081538)

# Transposed convolution

- Available in pytorch:  
<https://pytorch.org/docs/stable/generated/torch.nn.ConvTranspose2d.html>
- An excellent guide on this topic is here:  
<https://arxiv.org/abs/1603.07285v1>



# Challenge

Anders created a 50k/10k train/test split of the FaceScrub dataset. See next slide for download instructions.  
(full dataset here: <http://vintage.winklerbros.net/facescrub.html>)

- Denoising AutoEncoder (DAE) on FaceScrub:
  - Implement a convolutional AE to reconstruct face images
  - Now do the same but corrupt the images and teach the model to reconstruct the original
  - (Might be easier to train an over complete DAE)
- Bonus: Train a Variational AE (VAE) with a 2D latent space on MNIST
  - Plot the latent embeddings of the mnist images and compare with Figure 3B from the science paper: <https://www.cs.toronto.edu/~hinton/science.pdf>
  - Plot the corresponding grid of reconstructions

# FaceScrub

Drive: <https://drive.google.com/file/d/1Uzgc9c0MSYP4y9ia7mvwCGggjhTDEmrk/view?usp=sharing>

Download through python:

```
import numpy as np
from google_drive_downloader import GoogleDriveDownloader as gdd

gdd.download_file_from_google_drive('1Uzgc9c0MSYP4y9ia7mvwCGggjhTDEmrk', '~/img_align_celeba_50k.npz')
blob = np.load('~/img_align_celeba_50k.npz')
x_train, x_test = blob['x_train'], blob['x_test']
```